


# FLAPPY-LEARN

---

*genetic algorithm learns to fly*



# FLAPPY BIRD


.....

- released in 2013
- notoriously difficult
- addictive (*caused major distractions among 1<sup>st</sup> year FMI UB students*)
- **objective:** survive as long as possible
- **controls:** tap to "jump"
- **rules:** die upon hitting a pipe or leaving the screen

**APPROACH**

# THE BIRD'S BRAIN

.....



# THE BIRD'S BRAIN

---


- simple neural network
- inputs: horizontal & vertical distance to next gap
  - distance to screen edge "contained" in  $y$  component
  - normalized  $[-1, +1]$
- hidden layer: 5 neurons
  - *tanh* non-linearity (we don't use backprop)
- output: % willingness to jump
  - *sigmoid* activation (it's a brain, right?)
  - fires if over 0.5

# THE WORLD

---

- barrier position: anywhere from top to bottom edge
- gaps get increasingly smaller
  - until birds can no longer fit  $\Rightarrow$  force next generation
  - decay exponentially
- a bird's brain is simulated multiple times per second

# A GENERATION TAKING OFF



# GENETIC OPERATORS



# GENETIC SETTING

---

- genome: a bird's brain
- population: 10 genomes (one generation)
- fitness: the bird's age
- first generation: random network
  - weights:  $\sim U[-1, +1]$
  - biases: zero
- next generation: *select, crossover and mutation*
  - if none managed to pass a single barrier, wipe out (but keep best one)

# SELECTION

---


- *who advances to the next generation*
- fitness-proportionate roulette wheel selection
- keep the 3 fittest (elitism)
- pick 2 lucky ones, at random
- add 1 random genome
  - combat staleness

# HAVING FAITH PAYS OFF

Gen 10

Max 18

2/10 Alive



fittest

lucky

random

crossover

# CROSSOVER

---

- *offspring of two parents (one dominant, one recessive)*
- uniform recombination, binary mask
- weights: inherited from dominant
- biases: 25% from recessive
- 3 couples:
  - fittest two
  - two random winners
  - two randomly selected

# OVERCOMES ITS PARENTS

Gen 22

Max 30

3/10 Alive



fittest

lucky

random

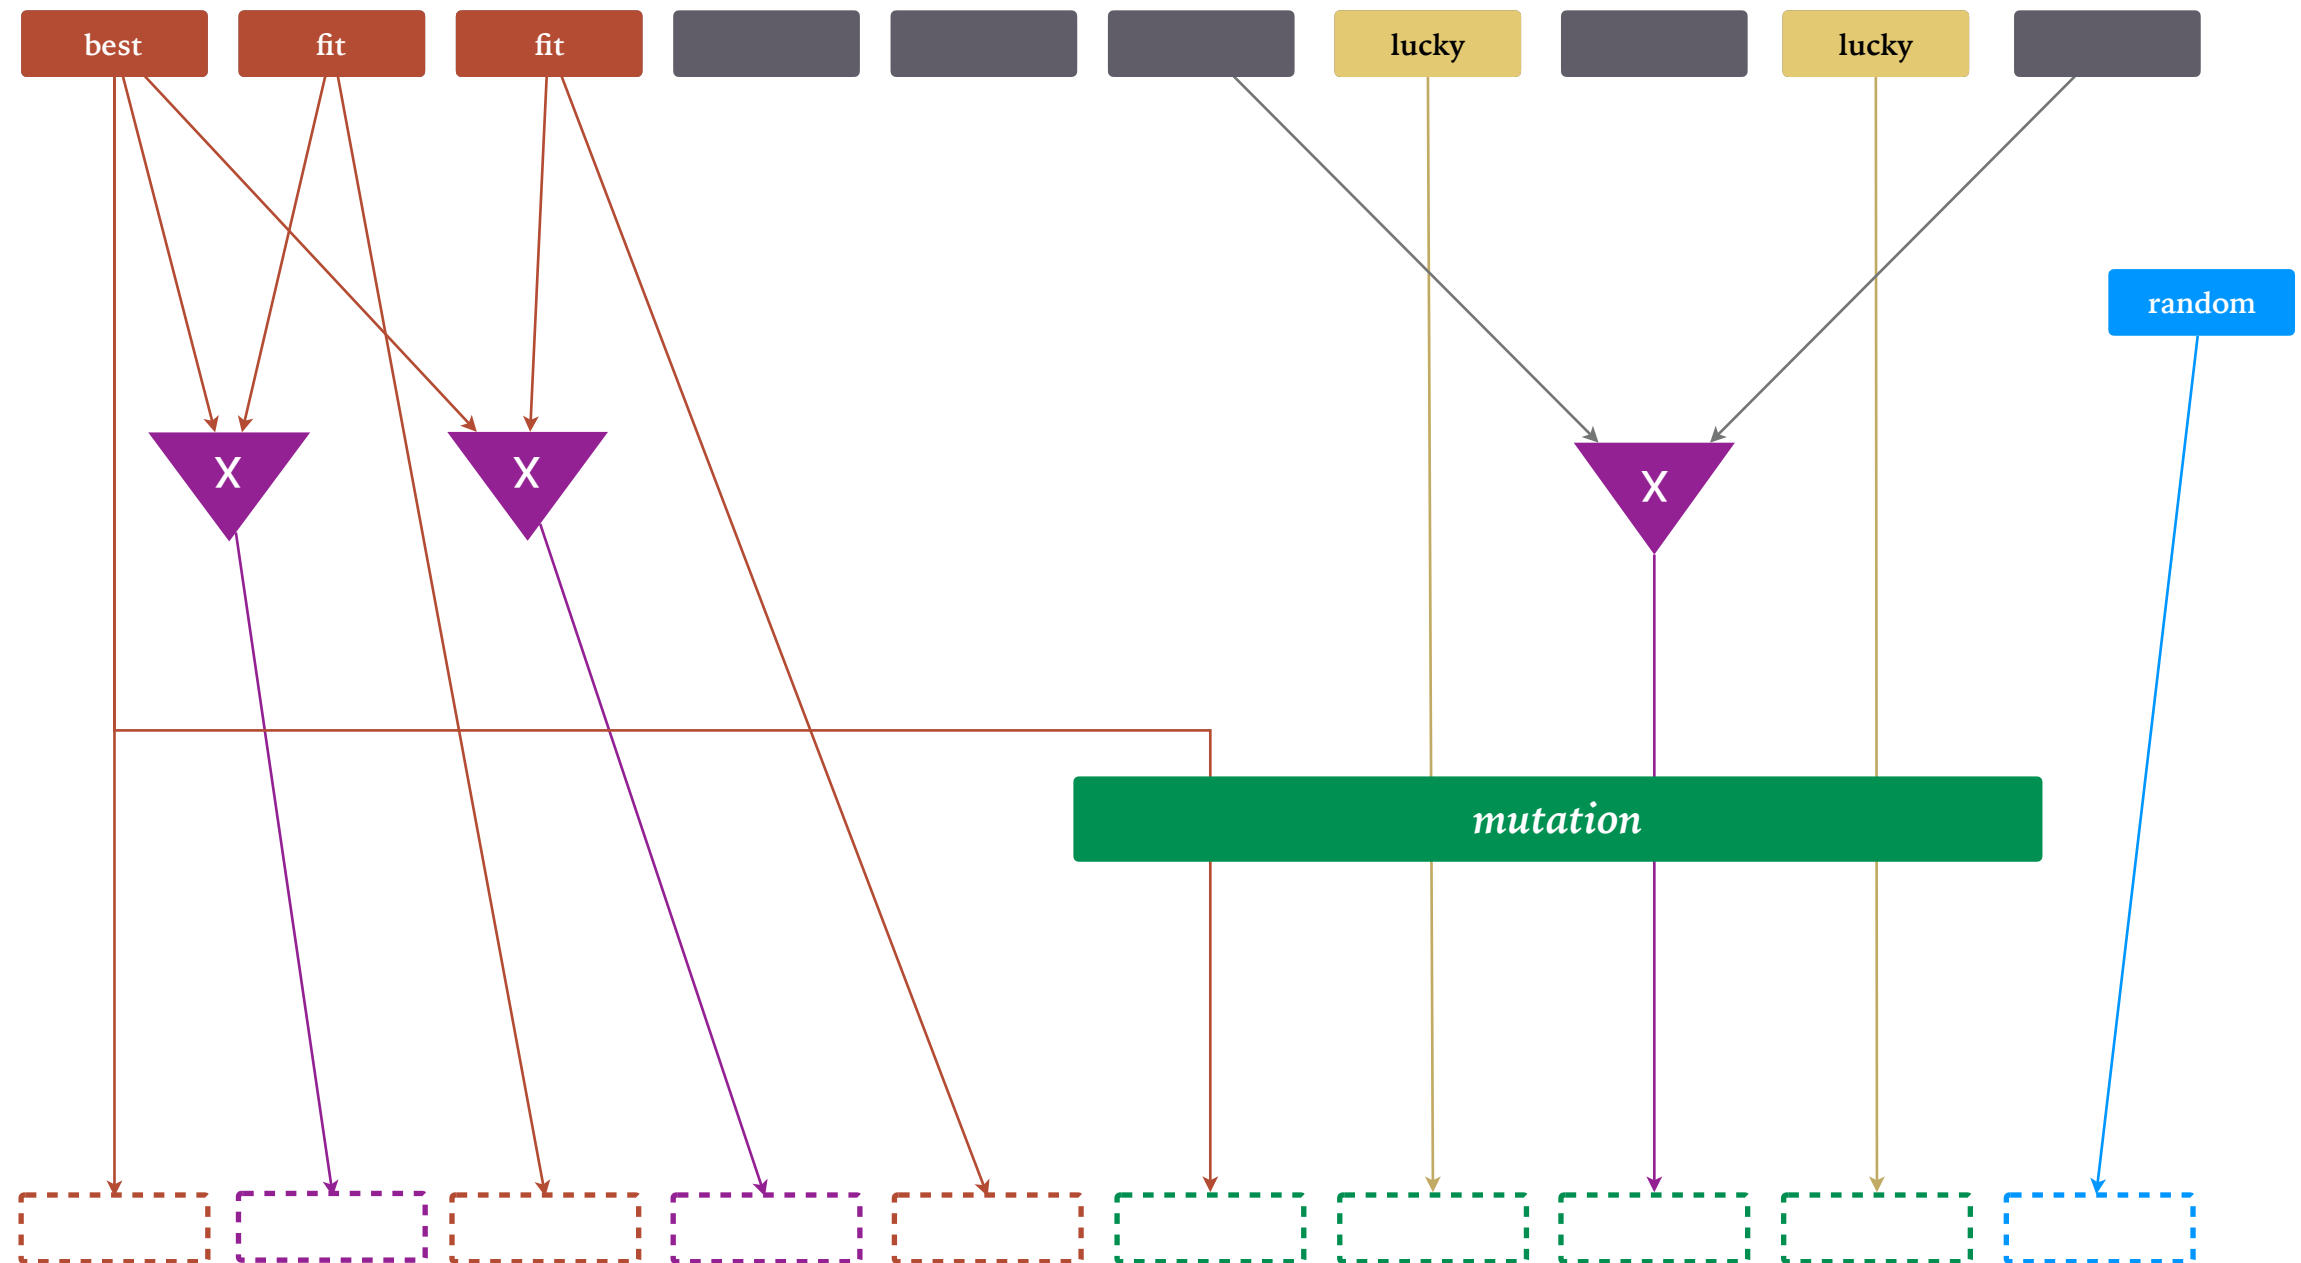
crossover

# MUTATION

---

- *vary genomes slightly*
- real-valued, uniform mutation
- 20% chance to apply to each neuron's weight
- $w \mapsto aw + b; \quad a \sim U[0.85, 1.15], \quad b \sim U[-0.2, +0.2]$
- mutants:
  - the best one
  - the lucky ones
  - offsprings

*current generation*



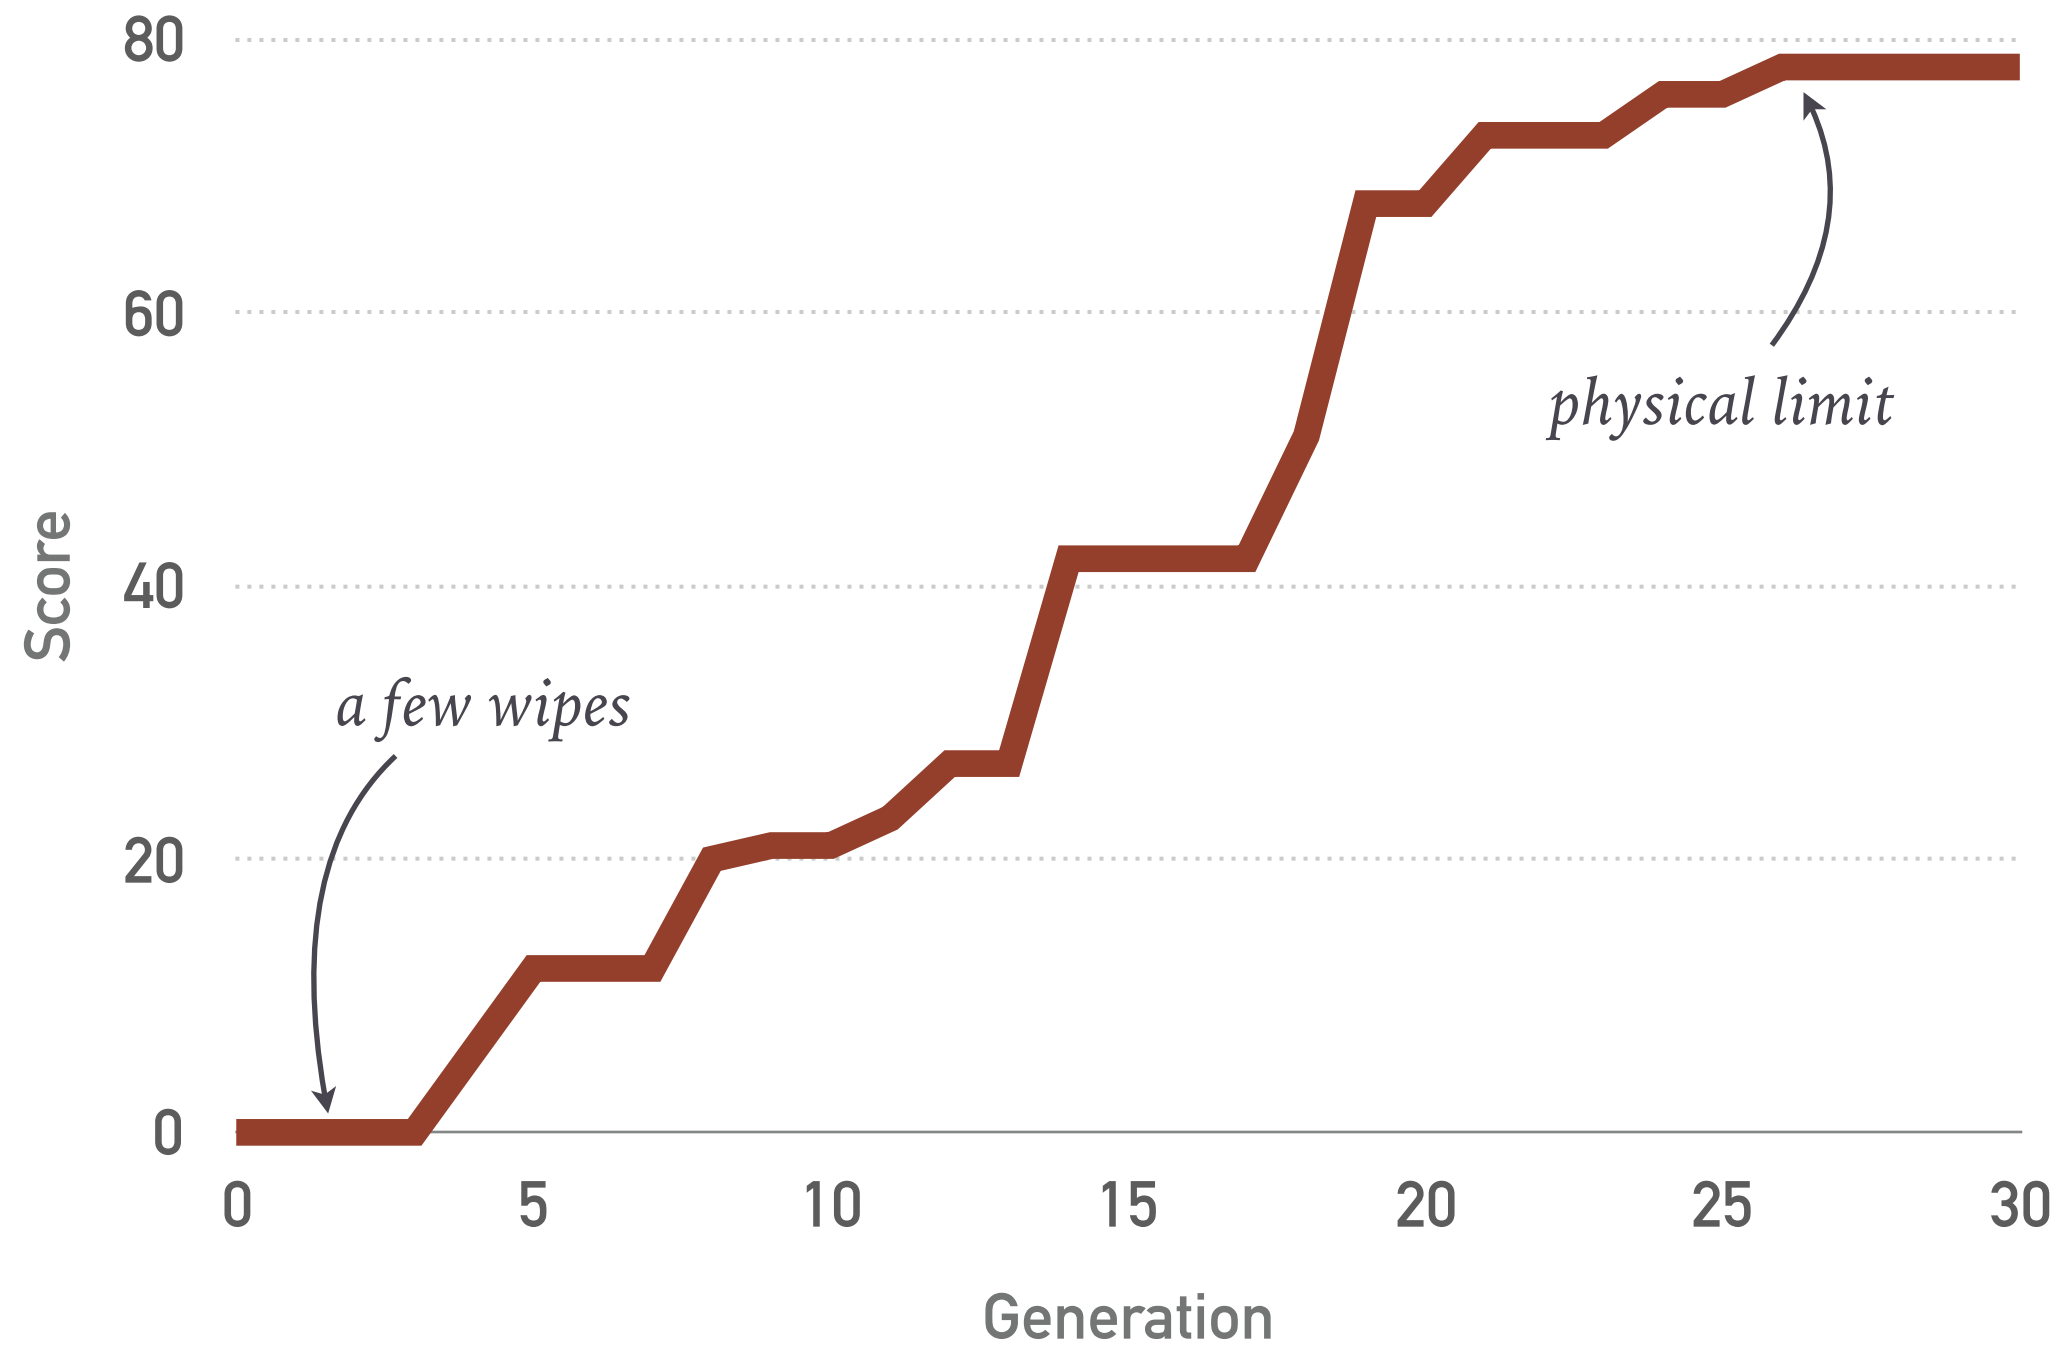
*next generation*

# RESULTS



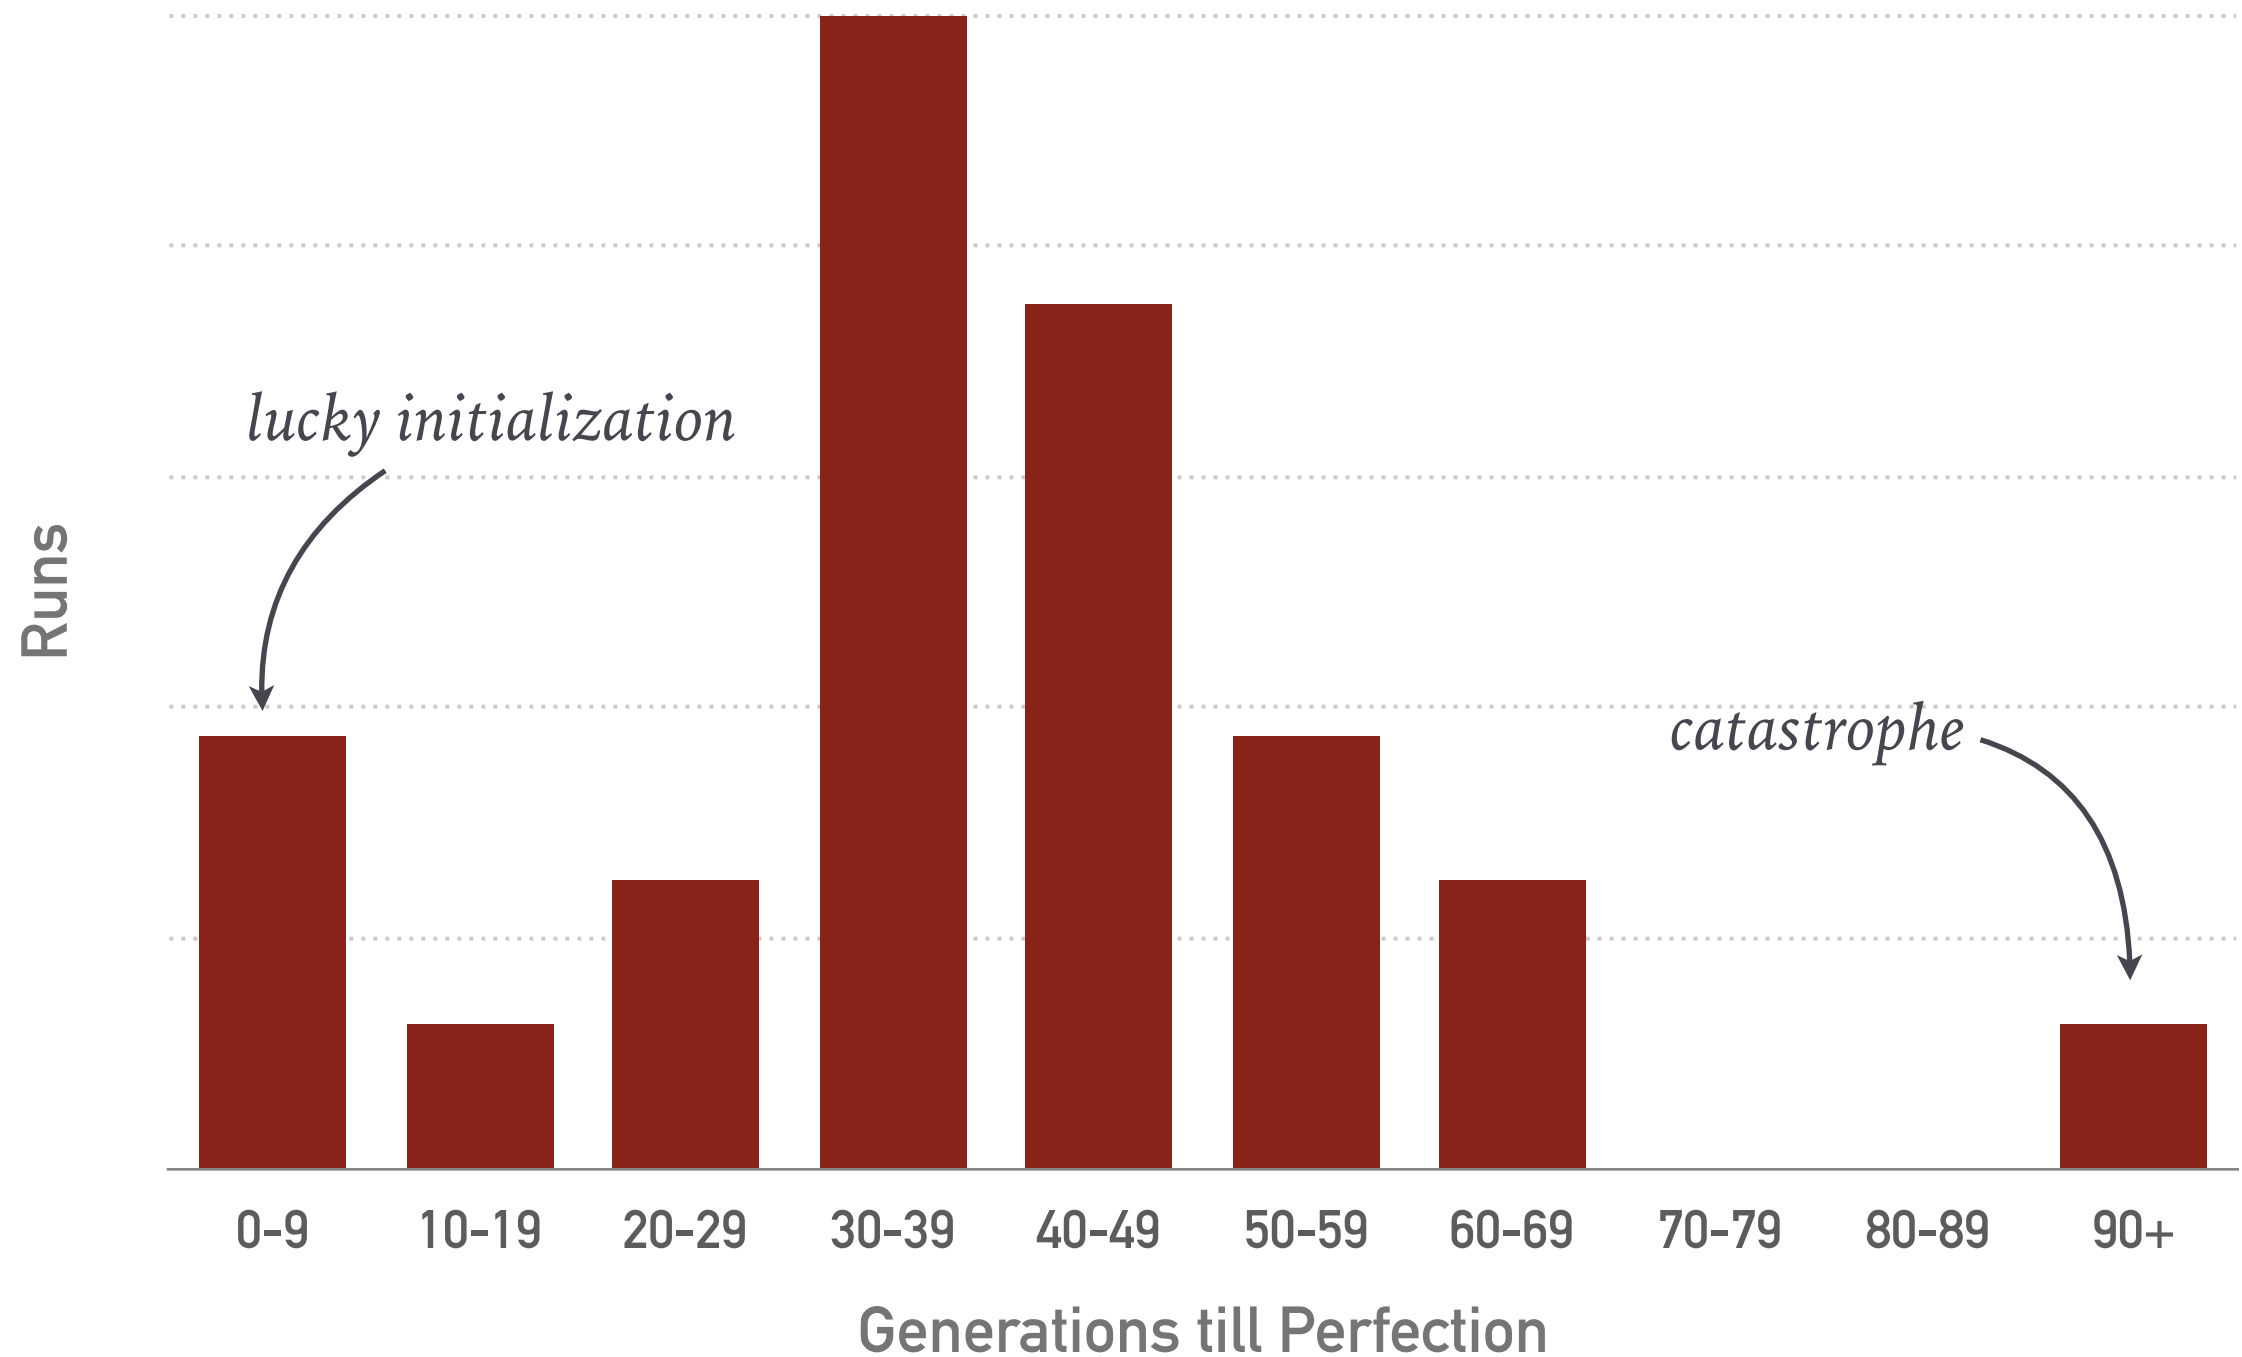
# LEARNING SPEED

---



# CONVERGENCE SPEED

---



# A CHAMPION IS BORN



# TRY IT YOURSELF!

---

- online demo: [stefann.eu/flappy-learn](http://stefann.eu/flappy-learn)
- checkpointing, can save/load the bird's brain
  - repo contains parameters of the champion
- source: [github.com/stefan-niculae/flappy-learn](https://github.com/stefan-niculae/flappy-learn)
- written in JavaScript from the ground up
  - using the Phaser game engine
  - assets from [kenney.nl](http://kenney.nl) (paid)

**THANK YOU!**